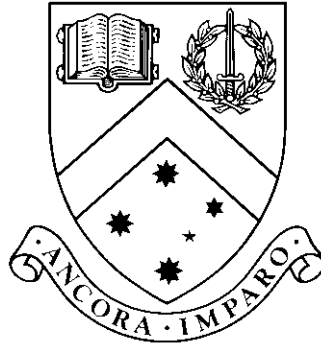


School of Computer Science (BICA)
Monash University



Literature Review, 2017

Review of optimal multi-agent Pathfinding algorithms and usage in warehouse automation

Phillip Wong

Supervisors: Daniel Harabor,
Pierre Le Bodic

Contents

1	Introduction	1
1.1	Problem Overview	1
1.2	Use of Mixed Integer Programming	1
2	Related Work	1
3	Algorithm	1
3.1	High Level	2
3.2	Path Generation	6
3.3	Path Assignment	6
3.4	Temporal A* with collision penalties	7
4	Resolving conflicts	7
5	Master problem formulation	8

1 Introduction

In our multi-agent pathfinding problem, we have an environment containing a set of k agents on a 4 directional grid-map. Each agent aims to find a path to their goal such that no agent will collide with another agent at any time.

Hence we have a centralized agent coordinator which aims to resolve path collisions.

1.1 Problem Overview

Inputs:

- Grid map
- Set of agents with
 - Start location
 - Goal location

Outputs:

- Collision free path from the start to goal location for each each agent

INSERT IMAGE OF EXAMPLE PROBLEM

1.2 Use of Mixed Integer Programming

Branch and bound

Column generation

Branch and price

2 Related Work

CBS

ICTS

Centralised A*

That Network Flow paper

3 Algorithm

Each agent knows of a number of paths to the goal. Using a Mixed Integer Program (MIP) we are able to solve the assignment problem of choosing a combination of paths such that no assigned two paths has a collision. Hence the two main parts of the algorithm:

Any specifics about the MIP I should mention here?

1. **Path Assignment** The Mixed Integer Program, Responsible for choosing a assigning paths to agents, while finding a conflict-free solution.
2. **Path Generation:** Responsible for generating alternative paths for agents whose current paths are in collision.

3.1 High Level

At the high-level we describe the cycle of path generation and assignment.

Initializing: Starting on line ??, we find the shortest path to each agent's goal and add it to their path bank. A path bank is simply a pool of valid paths to each agent's goal, these paths can be non-optimal.

Calling the MIP: Starting the loop, we call AssignPathsMIP. Here we create and run the Mixed Integer Program. This method will choose suitable paths from each agent's path pool. Importantly we have not yet described the second parameter, pathCollisions. Hence we do not know if the assignment is collision-free.

Checking collisions: So we must check if the solution the MIP gave us is valid.

No collision?: If there was no collision, then the solution the MIP assigned us was valid and we can terminate the algorithm.

Collision occurred: If there was a collision, we must perform extra steps. The goal here is to penalize this collision and generate new alternative paths. Hence we add this collision to the pathCollisions (used in the AssignPathMIP method), this later creates a hard constraint saying path i for agent n and path j for agent m cannot exist at the same time. Hence the MIP will not be able to choose these two in combination.

Choosing an agent: This method deals with choosing the next agent who we generate a path for. Here we choose the agent with the smallest Δ . This value describes the length between the shortest path and the current path. By choosing the agent with the smallest Δ we can ensure optimality as the makespan of the solution grows incrementally *IMPORTANT: describe this better.*

Penalties: Now we penalize the collision. We do this by looking at the perspective of the chosen agent. So for a collision between $a_n p_i$ and $a_m p_j$, say we have chosen a_n . Hence we are interested in penalizing against p_j . The penalty applies a negative value to any action that could lead to a collision against p_j **EXAMPLE HERE OR LINK TO EXAMPLE?**

Looping: Now we repeat this iteration of assignment and applying penalties. By ensuring that the makespan of the solution space grows incrementally by 1. We can generate an optimal solution, as the MIP will find a possible combination of paths from each agent's path bank.

Procedure 1 high-level of the Algorithm

Input: *agents* : list of agents

Output: *void*

```

1: pathCollisions  $\leftarrow \{\}$ 
2: penalties  $\leftarrow \text{Map} < \text{Agent}, \text{Penalties} > ()$ 
3:
4: for agent : agents do
5:   agent.pathBank.add(ShortestPath(agent))
6:
7: loop
8:   assignedPaths  $\leftarrow \text{AssignPathsMIP}(\text{agents}, \text{pathCollisions})$ 
9:   collision  $\leftarrow \text{CheckForACollision}(\text{assignedPaths})$ 
10:  if collision is not none then
11:    pathCollisions.add(collision)
12:    chosenAgent  $\leftarrow \text{ChooseAgent}(\text{collision})$ 
13:    UpdatePenalties(chosenAgent)
14:    chosenAgent.pathBank.add(GenerateNextBestPath(agent, penalties))
15:  else
16:    break
```

Simple example To get a better intuition of the algorithm, here we step through a simple example.

	0	
1	2	3
4	5	6

Table 1: A simple MAPF problem. $a_1 : (1, 6), a_2 : (0, 3)$

1. **Initializing:** Generate and store the shortest path in each agent's path bank. Since this is first iteration, we will be finding the shortest path for **all** agents. In later iterations, we only generate a path for **one** agent at a time.
 - $a_1 p_1 : \{1, 2, 3, 6\}$
 - $a_2 p_1 : \{0, 2, 3\}$
2. **Assign paths MIP:** Create a MIP from the agent's path bank and path constraints (currently none) and use the solution to assign paths to agents
 - $a_1 assigned : \{1, 2, 3, 6\}$
 - $a_2 assigned : \{0, 2, 3\}$
3. **Check for a collision:** Here we check the MIP path assignment by finding the first collision (if one exists). There is a path collision at timestep 2 where a_1 and a_2 are both on tile 2.
4. **Column generation? and constraints:** Since a collision exists, we now apply a constraint to the MIP describing that $a_1 p_1$ and $a_2 p_1$ cannot exist at the same time.
5. **Penalties:** In order to generate the next best path, we apply penalties incrementally when a collision occurs. The collision occurs between $a_1 p_1$ and $a_1 p_2$. Here we decide to generate a path for agent 1 so we penalize all of $a_1 p_2$, that is the tile of the path at timestep as so:
 - time 0, tile 0, penalty 1
 - time 1, tile 2, penalty 1
 - time 2, tile 3, penalty 1
6. **Generate next best path:** Here we generate and store the next best path. Unlike step 1, we now have penalties applied to certain tiles which lead our pathfinder to avoid collisions. Additionally unlike step 1, we choose one agent to generate a path for (here we have chosen agent 1).
 - $a_1 p_1 : \{1, 2, 3, 6\}$
 - $a_1 p_2 : \{1, 4, 5, 6\}$
 - $a_2 p_1 : \{0, 2, 3\}$
7. **Assign Paths MIP:** Create a MIP from the agent's path bank and use the solution to assign paths to agents. Here the MIP has now assigned $a_1 p_2$.
 - $a_1 assigned : \{1, 4, 5, 6\}$
 - $a_2 assigned : \{0, 2, 3\}$
8. **Check for a collision:** Here we check the MIP path assignment by finding the first collision (if one exists). There is no collision so the loop terminates and we are satisfied that this path assignment has no collisions.

	0	
1	2	3

Table 2: A medium complexity MAPF problem. $a_1 : (1, 3), a_2 : (3, 1)$

1. **Generate next best path:**

- $a_1p_1 : \{1, 2, 3\}$
- $a_2p_1 : \{3, 2, 1\}$

2. **Assign paths MIP:** Create a MIP from the agent's path bank and path constraints (currently none) and use the solution to assign paths to agents

- $a_1assigned : \{1, 2, 3\}$
- $a_2assigned : \{3, 2, 1\}$

3. **Check for a collision:** Here we check the MIP path assignment by finding the first collision (if one exists).

- collision at time 2 : a_1p_1, a_2p_2

4. **Path Collision:** Since a collision exists, we now apply a constraint to the MIP describing that a_1p_1 and a_2p_1 cannot exist at the same time.

5. **Penalties:** In order to generate the next best path, we apply penalties incrementally when a collision occurs. The collision occurs between a_1p_1 and a_1p_2 . Here we decide to generate a path for agent 1 so we penalize all of a_1p_2 , that is the tile of the path at timestep as so:

- time 0, (any to tile 3), penalty 1
- time 1, (any to tile 2), penalty 1
- time 2, (any to tile 1), penalty 1

6. **Generate next best path:** Here we generate and store the next best path. Unlike step 1, we now have penalties applied to certain tiles which lead our pathfinder to avoid collisions. Additionally unlike step 1, we choose one agent to generate a path for (here we have chosen agent 1).

- $a_1p_1 : \{1, 2, 3\}$
- $a_1p_2 : \{1, 1, 2, 3\}$
- $a_2p_1 : \{3, 2, 1\}$

7. **Assign Paths MIP:** Create a MIP from the agent's path bank and use the solution to assign paths to agents. Here the MIP has now assigned a_1p_2 .

- $a_1assigned : \{1, 1, 2, 3\}$
- $a_2assigned : \{3, 2, 1\}$

8. **Check for a collision:** Here we check the MIP path assignment by finding the first collision (if one exists).

- collision at time 2 : a_1p_2, a_2p_1

9. **Path Collision:** Since a collision exists, we now apply a constraint to the MIP describing that a_1p_2 and a_2p_1 cannot exist at the same time.

10. **Penalties:** a_1p_2 and a_2p_1 . Here we decide to generate a path for agent 2 so we penalize all of a_1p_2 .
 - time 0, (any to tile 1), penalty 1
 - time 1, (any to tile 1), penalty 1
 - time 2, (any to tile 2), penalty 1
 - time 3, (any to tile 3), penalty 1
 - time 2, (tile 2 to tile 1), penalty 1
11. **Generate next best path:** Here we generate and store the next best path for agent 2.
 - $a_1p_1 : \{1, 2, 3\}$
 - $a_1p_2 : \{1, 1, 2, 3\}$
 - $a_2p_1 : \{3, 2, 1\}$
 - $a_2p_2 : \{3, 3, 2, 1\}$
12. **Assign Paths MIP:** Create a MIP from the agent's path bank and use the solution to assign paths to agents.
 - $a_1assigned : \{1, 2, 3\}$
 - $a_2assigned : \{3, 3, 2, 1\}$
13. **Check for a collision:** Here we check the MIP path assignment by finding the first collision (if one exists).
 - collision at time 2 : a_1p_1, a_2p_2
14. **Path Collision:** Since a collision exists, we now apply a constraint to the MIP describing that a_1p_1 and a_2p_2 cannot exist at the same time.
15. **Penalties:** a_1p_1 and a_2p_2 . Here we decide to generate a path for agent 1 so we penalize all of a_1p_2 .
 - time 0, (any to tile 3), penalty 2*
 - time 1, (any to tile 2), penalty 1*
 - time 1, (any to tile 3), penalty 1
 - time 2, (any to tile 1), penalty 1*
 - time 2, (tile 2 to tile 3), penalty 1
 - time 2, (any to tile 2), penalty 1
 - time 3, (any to tile 1), penalty 1
16. **Generate next best path:** Here we generate and store the next best path for agent 2.
 - $a_1p_1 : \{1, 2, 3\}$
 - $a_1p_2 : \{1, 1, 2, 3\}$
 - $a_1p_2 : \{1, 1, 2, 2, 3\}$
 - $a_2p_1 : \{3, 2, 1\}$
 - $a_2p_2 : \{3, 3, 2, 1\}$

3.2 Path Generation

This method aims to generate paths for an agent. The first iteration of the path generation will use A* to quickly find the shortest path to the goal. This path will likely be in conflict. Next we use Temporal A* which aims to avoid collisions by applying penalties to tiles which are in conflict. We increment the penalties by 1 and iteratively finds paths of the next cost.

A generated path may have a longer length than the optimal single-agent solution. Hence we describe this difference to the optimal path as (Δ) . We do not generate a path for this agent until other agents in the deadlock have the same Δ . In this way we find an optimal solution.

However our implementation of Temporal A* is not complete. There are instances which it is unable to find a solution. If the optimal shortest path requires the agent to move past the goal tile. Then Temporal A* fails and will never expand past the goal. An example of this is can be seen in Figure 3. Hence in this algorithm we determine that agents are in **deadlock** when they have $n = 100$ number of collisions. If this occurs we start finding paths with Centralized A*.

$a1$		
$t2$		$t1$
		$a2$

Table 3: A deadlock occurs here. The optimal solution is

$a1$: d, r, r, u, d, w

$a2$: w, w, w, u, l, l

This deadlock occurs as the solution for $a1$ requires the algorithm to search past the goal.

Procedure 2 GenerateNextBestPath

Input: $agent, penalties$: tile penalties to be used for the TemporalAStar pathfinder.

Output: $path$: the generated path

```

1:  $pathCollisions \leftarrow \{\}$ 
2: for agent  $a$  in  $\vec{a}$  do
3:    $path \leftarrow \{\}$ 
4:   if agent.pathBank is empty then
5:      $path \leftarrow AStar$ 
6:   else if  $\exists$  agents in deadlock with  $a$  then
7:      $path \leftarrow CentralizedAStar(agents\ in\ deadlock)$ 
8:   else
9:      $path \leftarrow TemporalAStar(agent, penalties)$ 
10: return  $path$ 

```

3.3 Path Assignment

At the core, this method simply calls the Master problem (5). The objective here is to assign paths to agents which look suitable. This solution may be invalid but will be checked later. These agents are assigned the penalty variable by the mip and are returned as output of this method.

Procedure 3 AssignPaths

Input: \vec{a} : list of agents \vec{c} : list of collisions

Output: \vec{p} : list of agents who were assigned the penalty

- 1: $solution \leftarrow RunMasterProblem(\vec{a}, \vec{c})$
 - 2: **for** a in $solution.assignedAgents$ **do**
 - 3: assign collision-free path to a as described by $solution$
 - 4: **return** $solution.penaltyAgents$
-

3.4 Temporal A* with collision penalties

Here we describe our time-expanded A* implementation. The goal of this algorithm is to incrementally find paths of increasing cost to ensure optimality. We do this applying a penalty to tiles when collisions occur. Collision detection is described in Section 3.2.

Procedure 4 Temporal A* Node

Input: $parent, tile, goal, penalty$

Output: $Node$: a new Temporal A* Node

- 1: $h \leftarrow Heuristic(tile, goal)$
 - 2: **if** $parent$ is *null* **then**
 - 3: $time \leftarrow 0$
 - 4: $g \leftarrow 0$
 - 5: **else**
 - 6: $time \leftarrow parent.time + 1$
 - 7: $g \leftarrow parent.cost + 1$
 - 8: $f \leftarrow g + h + penalty$
 - 9: **return** $this$
-

Procedure 5 Temporal A* comparator

Input: node a , node b

Output: *bool*

- 1: **if** $a.f$ equals $b.f$ **then**
 - 2: **if** $a.penalty$ equals $b.penalty$ **then**
 - 3: **return** $a.g < b.g$
 - 4: **else**
 - 5: **return** $a.penalty < b.penalty$
 - 6: **else**
 - 7: **return** $a.f < b.f$
-

The temporal A* algorithm aims to find a path from start to goal. Additionally our variant takes in a set of collision penalties. A collision penalty describes that the agent should try to avoid moving from tile a to tile b at timestep t . The collision penalty is used to calculate the f value for each node and determines the priority for expansion.

struct {}

4 Resolving conflicts

1. Given a set of paths, S which contains all agent's path, find a new path for each agent their goal and add it to S
2. Detect any path collision for each path

Procedure 6 Temporal A* with collision penalties

Input: s : start, g : goal, \vec{c} : penalties, a vector of collision penalties

Output: p : path

```
1:  $path \leftarrow \{\}$ 
2: if  $s$  or  $g$  is not valid or  $s$  equals  $g$  then
3:   return  $path$ 
4:  $open \leftarrow PriorityQueue < Node > ()$ 
5: push  $Node(null, s, g, 0)$  into  $open$ 
6: while  $open$  is not empty do
7:    $current \leftarrow open.pop()$ 
8:   for  $action$  in  $\{up, down, left, right, wait\}$  do
9:      $nextTile \leftarrow$  perform  $action$  on  $current$ 
10:    if  $\exists node$  at time  $current.time + 1$  on tile  $nextTile$  then
11:      update  $node$  by comparing  $node.parent$  and  $current$ , taking the lowest  $f$ 
12:    else
13:       $penalty \leftarrow penalties[current.time + 1][current.tile][nextTile]$ 
14:       $newNode \leftarrow Node(current, nextTile, g, penalty)$ 
15: while  $current$  is not  $null$  do
16:   add  $current$  to the front of  $path$ 
17:    $current \leftarrow current.parent$ 
18: return  $path$ 
```

3. Convert the paths to MIP variables and path collisions to constraints
4. Repeat 1. if there is not a valid solution found i.e the optimal solution contains a path collision

5 Master problem formulation

Each agent is given *one or more* paths to their goal. The master problem aims to assign one path to every agent while minimizing the path distance and avoiding path collisions.

- **Agents:** The MIP is given a set of agents, A
- **Potential paths:** Every agent has a set of paths, P describing a unique path from an agent's position to their goal.
- **Penalty:** A penalty variable, q_i is added for every agent in the case that all the agent's paths are in collision. If the MIP chooses this penalty in the solution, then the MIP was unable to find an assignment of paths which resulted in a conflict-free solution. The cost of the penalty is set to be larger than the expected maximum path length (here it is 1000).
- **Conflicts:** A set of conflicts, C is provided to the MIP. A conflict is a set of paths, describing that these paths are not allowed to be assigned at the same time as they will cause a conflict.

We generate a variable for each path and the cost is set to the path length.

We specify an agent's path as p_{ij} . Penalty q_i . Path collision as c_{nm} .

$$\min \sum_{i \in A} \sum_{j \in P_i} (d_{ij} * p_{ij}) + q_i \quad (1)$$

$$\text{subject to } \sum_{j \in P_i} (p_{ij}) + q_i = 1, \forall i \in A \quad (2)$$

$$\sum_{m \in P_c} p_{cm} \leq 1, \forall c \in C \quad (3)$$

$$p_{ij} \in 0, 1, \forall i, j \quad (4)$$

$$q_i \in 0, 1, \forall i \quad (5)$$

For example ?, our generated variables are: $5a_0p_0 + 5a_0p_1 + 1000q_0 + 2a_1p_0 + 2a_1p_1 + 1000q_1$.

Agents are assigned

References

- De Koster, R., Le-Duc, T. and Roodbergen, K. J. (2007). Design and control of warehouse order picking: A literature review, *European Journal of Operational Research* **182**(2): 481–501.
- Wurman, P. R., D’Andrea, R. and Mountz, M. (2008). Coordinating hundreds of cooperative, autonomous vehicles in warehouses, *AI magazine* **29**(1): 9.